

# Crash Recovery in a Distributed Data Storage System

Lampson and Sturgis 1979

Presented by Raymond Xu

# Motivation

- Crashes may cause the state of a system to become inconsistent

```
transfer(src, dst, amount) {  
    src_bal = read(src)  
    write(src, src_bal - amount)  
    dst_bal = read(dst)  
    write(dst, dst_bal + amount)  
}
```

# Motivation

- Crashes may cause the state of a system to become inconsistent

```
transfer(src, dst, amount) {  
    src_bal = read(src)  
    write(src, src_bal - amount)  
    dst_bal = read(dst)  
    write(dst, dst_bal + amount)  
}
```

## Goal: Atomic transactions


# Foundation

- Distributed storage system
  - Clients and servers
  - Clients talk directly to servers containing data
  - Servers can send messages to each other
- Consistency
  - A state is consistent if it satisfies the invariant of the system
- Transaction
  - A sequence of read/write commands that, after recovery from a crash, will have all been executed or not executed at all

# Failure model

- Events can be **desired** or **undesired**
- Undesired events can be **expected (errors)** or **unexpected (disasters)**
- The paper's algorithm assumes no disasters but any number of errors

# Abstraction Lattice



# Physical System

# Disk storage

`Get(at: Address) returns (status: (good, looksBad), data: Dblock)`

`Put(at: Address, data: Dblock)`



# Disk storage

Get(at: Address) returns (status: (good, looksBad), data: Dblock)

- (error) Soft read error
  - page is good, but status is bad
- (disaster) Persistent read error
  - soft read error but successively
- (disaster) Undetected error
  - page is bad but status is good, or data returned is wrong block

# Disk storage

Put(at: Address, data: Dblock)

- (error) Null write
  - page is unchanged
- (error) Bad write
  - data is written to page, but page status becomes bad

# Disk storage

- **Decay set:** pages that may decay together
- **Decay:** spontaneous event that changes a subset of a decay set from good to bad
- *Assumption:* it is possible to partition disk storage into pairs of units that are not decay-related
- (error) Infrequent decay
  - Interval  $T_D$  separates all decays in a unit
- (error) Revival
  - A page goes from bad to good
- (disaster) Frequent decay
  - No interval  $T_D$  in play
- (disaster) Undetected error
  - A page's data changes

# Processor

- A crash is an error that resets the state of the processor (volatile state) to a standard value
- *Assumption:* No other processor errors are allowed — any malfunction can be detected and converted into a crash

# Communication

`Send(to: Processor, data: Mblock)`

`Receive returns (status: (good, bad), data: Mblock)`

# Communication

Send(to: Processor, data: Mblock)

Receive returns (status: (good, bad), data: Mblock)

- (error) Loss
  - some message is destroyed
- (error) Duplication
  - some new message identical to an existing one is created
- (error) Decay
  - some message status changes from good to bad
- (disaster) Undetected error
  - some message changes from bad to good, or the data or recipient process change for a good message

# Atomicity

An action is atomic if it is:

1. **Unitary:** If the action returns, it was carried out completely; if the system crashes before the action returns, then after the crash the action has been carried out completely or not started

*AND*

2. **Serializable:** When actions are done by concurrent processes, the result is as if the individual actions were carried out one at a time in some order. Action order must also be preserved within a process.

# Atomicity

An action is atomic if it is:

1. **Unitary:** If the action is interrupted by a system crash before it has been carried out completely; if the system recovers, the action has been carried out completely.
2. **Serializable:** If the action is interrupted by a system crash, the result is as if the individual actions were carried out one at a time in some order. Action order must also be preserved within a process.

“Unfortunately, we are unable to make all the actions in our various abstractions atomic” (9)



Stable System

# Careful Storage

- CarefulGet
  - Get until good, or after  $n$  tries
    - Eliminates soft read errors
- CarefulPut
  - Put then Get, repeat until success
    - Eliminates null writes and bad writes

# Stable Storage

Stable page comprised of an ordered-pair of careful disk pages (that are not decay-related) and a monitor

- **StableGet**
  - CarefulGet from first page, if bad then CarefulGet from second page
- **StablePut**
  - CarefulPut to first page then to second page
    - Atomic
- **Cleanup**
  - CarefulGet both pages:
    - If both good and same data, then do nothing
    - If one bad, then CarefulPut the good data to the bad address
    - If both good but different data, then CarefulPut one data to the other

# Stable Storage

- Run Cleanup at initialization, after each crash, and at least every  $T_D$
- *Invariant:* Both pages cannot be bad, and if both are good then they both have the most recent data, except during a StablePut
  - If first page decayed, then the other page cannot decay in  $T_D$  and cannot experience a bad write, so Cleanup will fix the bad page
  - If first page suffered a bad write, it will be fixed by CarefulPut or Cleanup

# Stable Processors

1. Can use disk storage to store state
2. Can use stable storage to store state for crash recovery
3. Can use stable storage to construct stable monitors
  - a. Monitor: data wrapper that uses locking to ensure mutual exclusion
  - b. Stable monitor: all data besides the lock is in stable storage

# Remote Procedures

- Use globally unique ids for messages, attainable via a stable monitor
- Periodic retries, deduplicate by id
- Receivers track largest seen id from each processor, ignore ids  $\leq$  largest
- If receiving a message from a processor that you have no record of, request that processor's current id

# Compatible Actions

Weaker substitute for serializability

A history of actions is compatible if there is a serial history, in which there is no concurrency and each action occurs in isolation

The serial history must agree with the actual history:

- Mapping from serial history actions to actual history actions
- Actions that complete before a strong action starts will appear to occur before the strong action
- Actions that begin after a strong action ends will appear to occur after the strong action

# Stable Sets and Compound Atomic Actions



# Stable Sets

StablePut is good and atomic, but we want to support arbitrarily large data

A Stable Set has a unique id and a set of records

Pool: maintain a set of pages in stable storage for use by stable sets

# Stable Sets

## Atomic operations

- `Create(i: ID)`
- `Insert(s, t: StableSet, new: Record)`
- `Replace(s, t: StableSet, old, new: Record)`
- `IsEmpty(s, StableSet)`
- `IsMember(s: StableSet, r: Record)`

## Non-atomic operations

- `Enumerate(s: StableSet, p: procedure)`
- `Erase(s: StableSet)`

# Stable Sets

- All actions on a stable set are compatible
- But using atomic operations while non-atomic operations are in progress has arbitrary resolution
- The pool can be ordered in a ring to reuse space
- For a wide stable set (a stable set that spans more than one processor) store metadata in each processor (1 root, potentially many leaves)

# Compound Atomic Actions

**procedure**  $A = \mathbf{begin}$  *Save*;  $R$ ; *Reset* **end**

- If  $R$  is an action compatible with a set of actions  $S$ , then  $A$  is an atomic action compatible with  $S$ :
  - If the procedure crashes before the *Save*, it's a no-op
  - If it crashes after the *Reset*,  $R$  has been done completely and will not be done again because the saved state has been erased
  - If it crashes between the *Save* and the *Reset*, it will resume after the *Save* and restart  $R$

# Transactions

2 phases:

1. Record intentions, after last action here the transaction is committed
2. Do the writes

If a crash occurs after commit but before all writes are complete, restart phase 2

- Recording intentions must be atomic
- Store intentions in a stable set, store locks in another stable set, committing the transaction and doing the writes is a compound atomic action

# Transactions

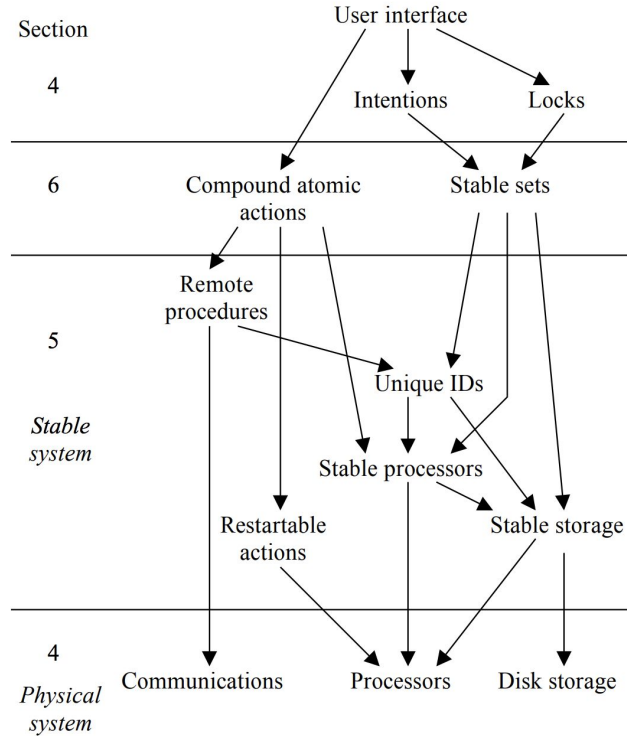
On each server store:

- Stable set containing Intentions for each page the server stores
  - Stable monitor protecting it
- Stable set containing transaction flags
  - Stable monitor protecting it
- If coordinator for a transaction, a root for a wide stable set containing transaction Intentions; a leaf if not coordinator but part of the transaction

# Transactions

- Client calls `Begin` on one of the servers, which becomes the coordinator
- Client calls `Read` and `Write` on the servers that contain the corresponding pages
- After all the `Write` calls return, client calls `End` or `Abort` on the coordinator

# Abstraction Lattice





# Crash Recovery in a Distributed Data Storage System

Lampson and Sturgis 1979

Presented by Raymond Xu